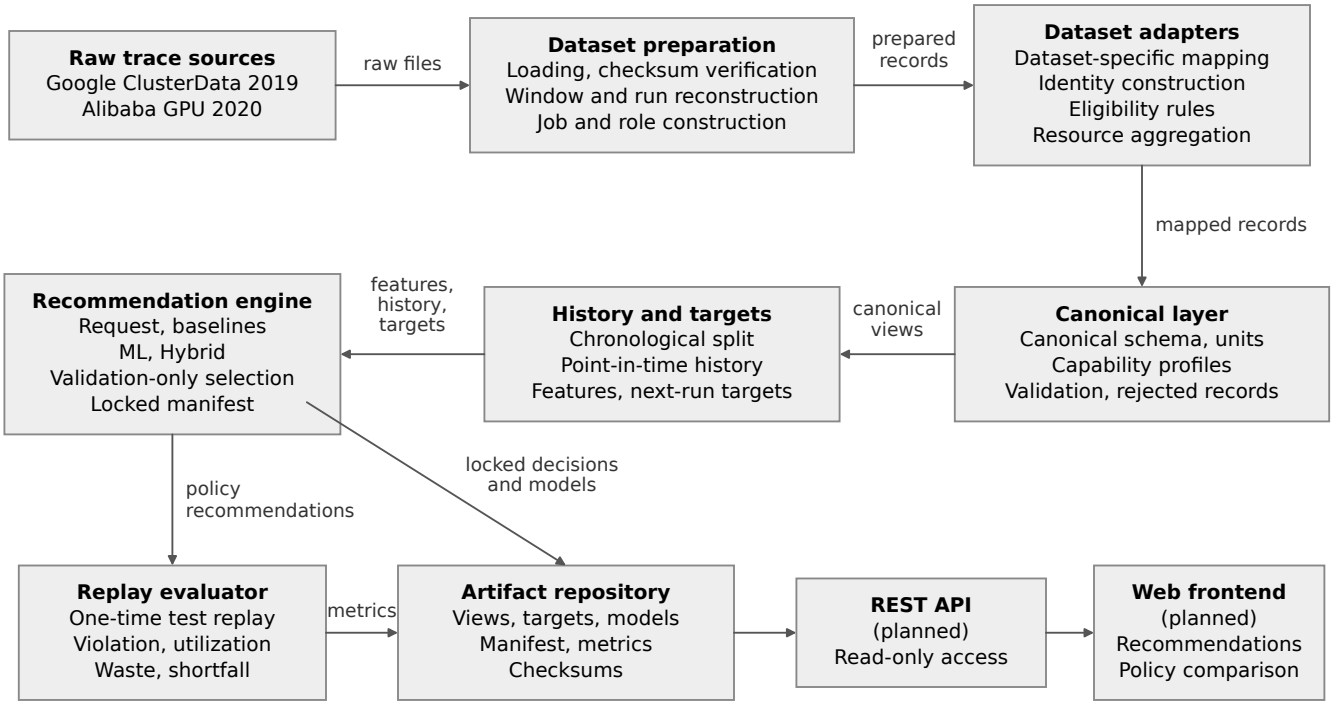
Raw trace sources
Google ClusterData 2019
Alibaba GPU 2020
raw files
Dataset preparation
Loading, checksum verification
Window and run reconstruction
Job and role construction
prepared records
Dataset adapters
Dataset-specific mapping
Identity construction
Eligibility rules
Resource aggregation
mapped records
Canonical layer
Canonical schema, units
Capability profiles
Validation, rejected records
canonical views
History and targets
Chronological split
Point-in-time history
Features, next-run targets
features, history, targets
Recommendation engine
Request, baselines
ML, Hybrid
Validation-only selection
Locked manifest
policy recommendations
locked decisions and models
Replay evaluator
One-time test replay
Violation, utilization
Waste, shortfall
metrics
Artifact repository
Views, targets, models
Manifest, metrics
Checksums
REST API
(planned)
Read-only access
Web frontend
(planned)
Recommendations
Policy comparison