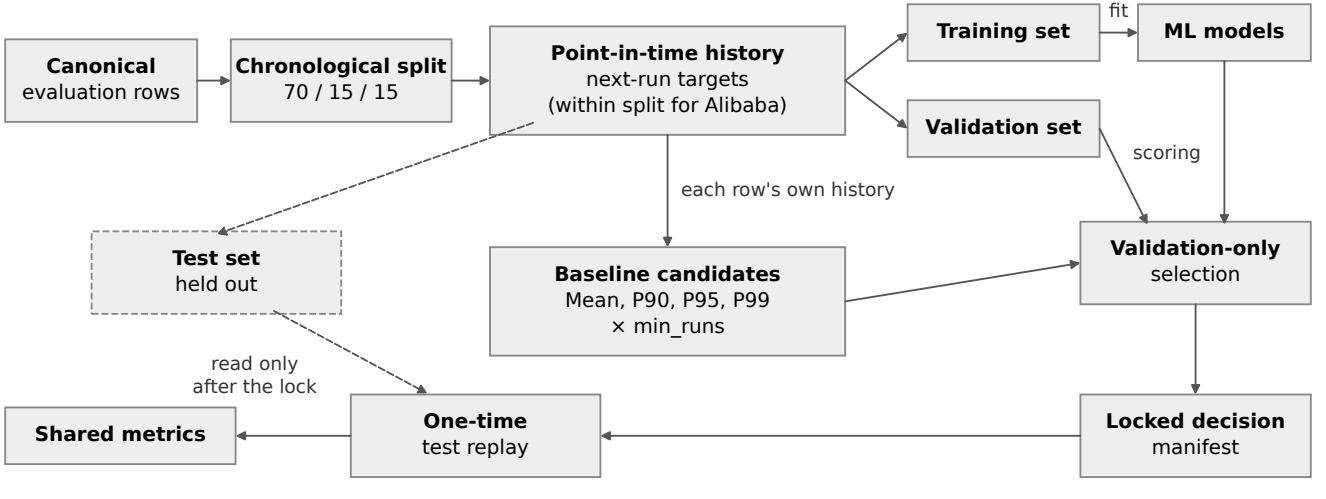
Canonical
evaluation rows
Chronological split
70 / 15 / 15
Point-in-time history
next-run targets
(within split for Alibaba)
Training set
fit
ML models
Validation set
scoring
each row's own history
Test set
held out
Baseline candidates
Mean, P90, P95, P99
× min_runs
Validation-only
selection
read only
after the lock
Shared metrics
One-time
test replay
Locked decision
manifest